

# Robust Adversarial Image Detection in Classification

Arti Jain, Grant Ovsepyan, Nithanth Ram

# Problem Statement

- Adversarial attacks on machine learning models can have catastrophic effects
  - Ex. Autonomous vehicle system trying to perform traffic sign classification
  - White-box attacks (access to model's gradients) and black-box attacks (no access to model)
- Investigate various conditioning techniques to explore the possibility of strengthening classification at inference time against white-box and black-box image perturbations

# Previous Work

## Color and Edge-Aware Adversarial Image Perturbations [1]

- New easy-to-compute subtle perturbations involving regions of smoothness and color contrast

## Adversarial Training for Free! [3]

- Add semi-supervised virtual adversarial noise to the input during training, acting as a regularizer

## n-ML: Mitigating Adversarial Examples via Ensembles of Topologically Manipulated Classifiers [2]

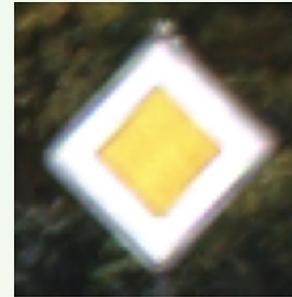
- Train multiple topologies so that ensembling these networks together provides more defense against adversarial examples

# Dataset

**Dataset:** GTSRB (German Traffic Sign Recognition Benchmark)

- 43 classes (speed limit, stop sign, pedestrian crossing, etc.)
- 39,210 training images

Original Image



## White Box Attacks

- FGSM (Fast Gradient Sign Method)
  - Fast and efficient method to generate attack
- PGD (Projected Gradient Descent)
  - Worst case boundaries of the model performance

FGSM



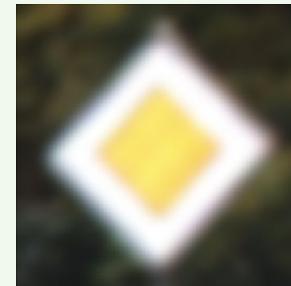
PGD



## Black Box Attack:

- Gaussian Blur
  - mimics real world scenario -> out of focus images

Gaussian Blur



# Our Approach

- Our approach was to experiment with image manipulation techniques to observe the change in misclassification from a perturbed dataset
  - Some of these processing techniques included utilizing canny edge maps, depth maps, contrast/light enhancement, de-blurring, de-noising
- We tried 3 methods:
  1. Concatenating canny edge maps to the input data and then passing it through our trained model
  2. Using different image processing techniques on the perturbed data and then running inference
  3. Ensembling separate classifiers/sub-neural networks

# Experiment Framework

- Fine-tune existing pre-trained image classifier models on traffic sign dataset
- Perform adversarial perturbations to dataset and measure degradation of model efficacy
- Tried mitigation technique
  - Concatenating Canny Edge image to original image
  - Unblurring/denoising image (MPRNet), contrast enhancement, K-means clustering, based on the color similarity
  - Ensembling and stacking models to get more robust classification
- Use F1-score and accuracy to track the performance across experiments

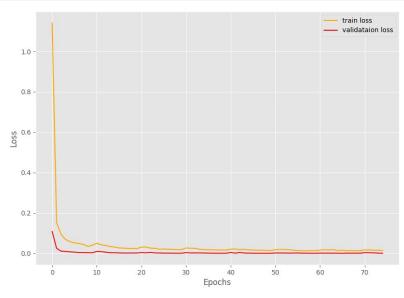
# Training

- Base Models: ResNet50, InceptionNetV3, MobileNetV3
  - ResNet50 fine-tuned ~ 76% test accuracy, 0.73 test F1-score
  - InceptionV3 fine-tuned ~ 68% test accuracy, 0.70 test F1-score
  - MobileNetV3 fine-tuned ~ 99.3% test accuracy, 0.99 F1-score
- Proceeding with MobileNetV3 as baseline fine-tuned model
- Apply adversarial perturbations with fixed parameters for control
  - Gaussian Blur: kernel\_size=5, sigma=2 (severe blur from visual inspection)
  - FGSM attack epsilon = 0.1, (1 iteration) (significant noise perturbation, additional noise would perturb image too much)
  - PGD attack epsilon = 0.1, alpha= 0.04, num\_iteration = 10 (less noise perturbation than FGSM, noise density between original unperturbed image and FGSM image)
- Decided to manually sort the classes for further inference using most readable images taken under reasonable light conditions
- Selected classes are: **0,1,2,11,12,14,16,28,34,35,39 (3810 images)**
- Baseline test accuracies on the adversarial images using MobileNetV3 on the selected classes are
  - Gaussian Blur: 89.61%
  - FGSM attack: 67.14%
  - PGD attack: 77.27%

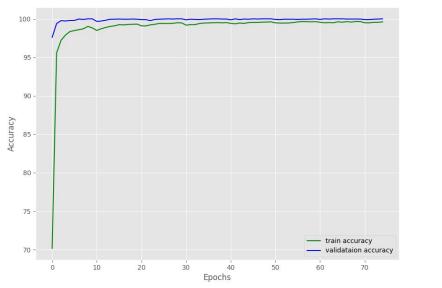
# Training and Validation Curves

## MobileNetV3:

training and validation loss

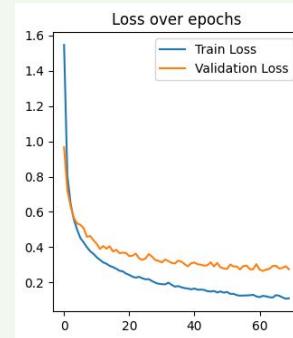


training and validation accuracies

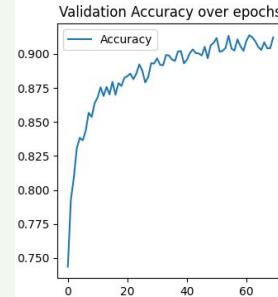


## ResNet50:

training and validation loss

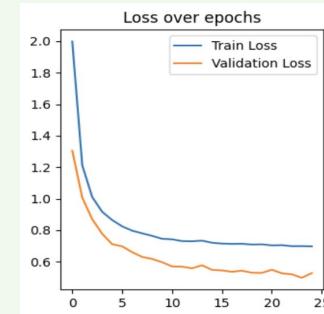


validation accuracy

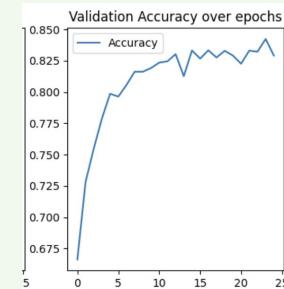


## InceptionNetV3:

training and validation loss

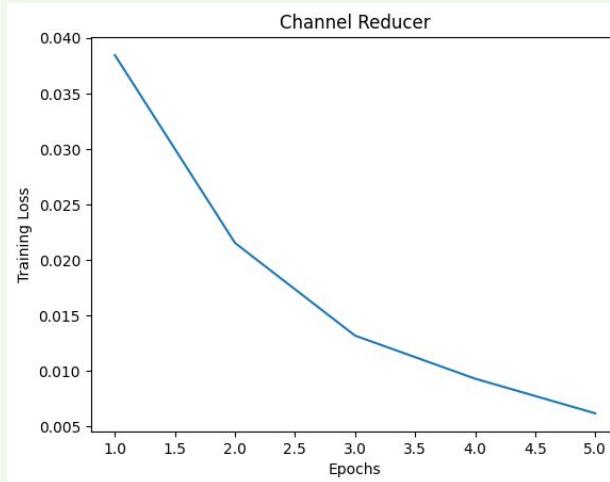


validation accuracy



# Solution Space (MobileNetV3)

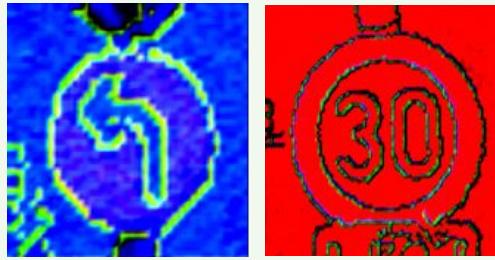
- Concatenate edge image and normal image before performing inference
  - a. Expand 3 channel image to 4 channel and pass through fine tuned model
  - b. See if model can't already judge augmented inputs with its current weights
- Trained a simple CNN (one layer) to convert 4 channels into 3
  - Used training data to create pairs of canny edge maps as the input, the original image as the label
  - 5 epochs, MSE Loss, LR = 0.001



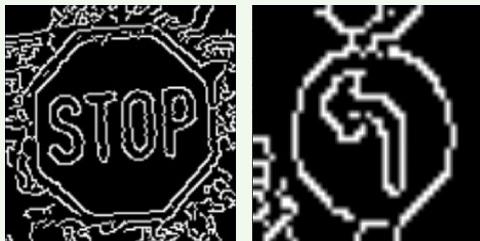
# Solution Space (MobileNetV3)



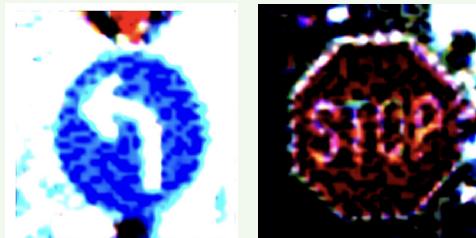
Perturbed Data with FGSM attack



Reducing from 4 channels to 3 channels  
with random CNN weights



Canny images on perturbed data



Reducing from 4 channels to 3 channels with  
trained CNN

# Solution Space (MobileNetV3)

|  | Test Accuracy | F1 Score |
|--|---------------|----------|
| Baseline with FGSM Attack              | 97.35%        | 0.986    |
| With Canny Edge and trained CNN        | 75.76%        | 0.82     |
| With Canny Edge and random CNN weights | 30.48%        | 0.29     |

|                                    | Test Accuracy | F1 Score |
|------------------------------------|---------------|----------|
| Baseline with Gaussian Blur Attack | 89.61%        | 0.92     |
| With Canny Edge and trained CNN    | 86.726%       | 0.897    |

|                                 | Test Accuracy | F1 Score |
|---------------------------------|---------------|----------|
| Baseline with PGD Attack        | 77.27%        | 0.84     |
| With Canny Edge and trained CNN | 81.42%        | 0.86     |

# Solution Space (MobileNetV3)

- 1) Denoising/Deblurring using MPRNet (off-shelf)
  - a) Intuitively, if FGSM and PGD add the noise to the image, then denoising those images should reverse noise addition. Similarly, deblurring should reverse the effects of Gaussian Blur

MPRNet was chosen because it utilizes multiple stages to progressively restore the image focusing on different aspects of restoration:

- 1) Encoder – Decoder Network: large receptive fields essential for understanding the overall image context.
- 2) Original Resolution Network: preserve the fine details necessary for high-quality restoration.
- 3) SAM: Supervised Attention Module (SAM) is used between each stage to refine the processed features is used to guide the attention mechanism
- 4) CSFF: Cross Stage Feature Fusion Technique, that ensures that integrates the features from different stages preventing the loss in details

- 2) Denoising/Deblurring using MPRNet + contrast enhancement
  - a) After inspecting some images, it was noted that some images might not have defined sign boundaries that blend in with environment, so contrast enhancement was introduced.

baseline\_original\_images



baseline\_gaussian\_blur



de-gaussian.blur



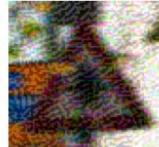
de-gaussian.blur\_light



baseline\_original\_images



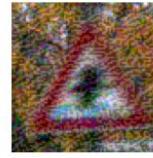
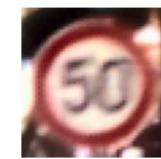
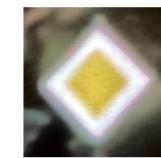
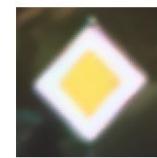
baseline\_fgsm\_attack



de-fgsm\_attack



de-fgsm\_attack\_light



baseline\_original\_images



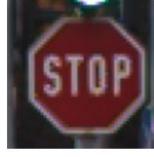
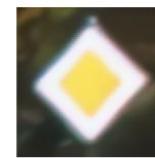
baseline\_pgd\_attack



de-pgd\_attack



de-pgd\_attack\_light



# Test accuracies on 3810 images on the selected classes (0,1,2,11,12,14,16,28,34,35,39)

| Image Set                                     | Test Accuracy | F1 Score |
|---|---------------|----------|
| Original Images                               | 99.30%        | 0.9953   |
| Gaussian Blur Baseline                        | 89.61%        | 0.9196   |
| Deblurred (MPRNet)                            | 87.74%        | 0.8998   |
| Deblurred (MPRNet) + Contrast enhancement     | 88.35%        | 0.8836   |
| FGSM Baseline                                 | 67.14%        | 0.7540   |
| Denoised FGSM (MPRNet)                        | 78.11%        | 0.8566   |
| Denoised FGSM (MPRNet) + Contrast enhancement | 82.23%        | 0.8737   |
| PGD Baseline                                  | 77.27%        | 0.8431   |
| Denoised PGD (MPRNet)                         | 83.76%        | 0.8953   |
| Denoised PGD (MPRNet) + Contrast enhancement  | 87.16%        | 0.9115   |

# Solution Space (MobileNetV3)

- 1) Manual clustering and pseudo-denoising based on k-Means clustering
  - a) Why?
    - i) From previous experiments the denoising approach improved the test accuracy on the PGD and FGSM images.
  - b) How?
    - i) Empirically determined most optimal k value = 5 that corresponds to number of most dominant colors in the image.
    - ii) Segment the images based on 5 dominant colors
    - iii) For each segmented group (not color, since there will be multiple groups for the same color):
      - 1) Average the pixel values on the adversarial image for each group
      - 2) Replace the values of that group with an averaged value

# Solution Space (MobileNetV3)

- 1) Color Averaging: pseudo-denoising based on k-Means clustering

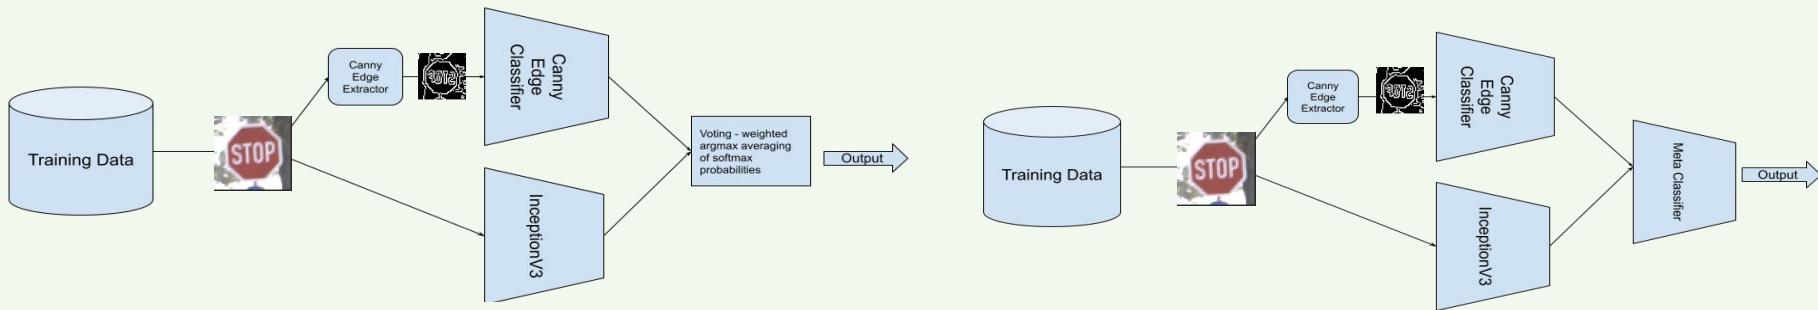


# Test accuracies on 3810 images on the selected classes (0,1,2,11,12,14,16,28,34,35,39)

| Image Set   | Test Accuracy | F1 Score |
|---|---------------|----------|
| Original Images   | 99.30%        | 0.9953   |
| Original Images + Color Averaging                               | 95.88%        | 0.9708   |
| Gaussian Blur Baseline + Color Averaging                        | 72.89%        | 0.7766   |
| Deblurred (MPRNet) + Color Averaging                            | 79.11%        | 0.8344   |
| Deblurred (MPRNet) + Contrast enhancement + Color Averaging     | 80.34%        | 0.8412   |
| FGSM Baseline + Color Averaging                                 | 69.19%        | 0.7677   |
| Denoised FGSM (MPRNet) + Color Averaging                        | 70.76%        | 0.7875   |
| Denoised FGSM (MPRNet) + Contrast enhancement + Color Averaging | 70.84%        | 0.7719   |
| PGD Baseline + Color Averaging                                  | 78.85%        | 0.8498   |
| Denoised PGD (MPRNet) + Color Averaging                         | 77.09%        | 0.8384   |
| Denoised PGD (MPRNet) + Contrast enhancement + Color Averaging  | 76.38%        | 0.8178   |

# Solution Space (InceptionV3)

- Stacking/Ensembling classifiers for edge image and normal image
  - Weighted voting - averaging the softmax probabilities from both and applying a weighting based on which classifier is better performant
  - Training a meta-classifier - overall neural network that has the independent image networks as sub networks



- Stacking/Ensembling failed - worse than random guessing ~ 45% accuracy
  - Canny Edge classifier had poor performance
  - Quality of edge extracted images for a lot of classes was poor - no edges detected

|             |  | Confusion Matrix |        |        |         |         |         |         |         |         |         |         |
|-------------|--|------------------|--------|--------|---------|---------|---------|---------|---------|---------|---------|---------|
|             |  | Predicted Labels |        |        |         |         |         |         |         |         |         |         |
|             |  | Class0           | Class1 | Class2 | Class11 | Class12 | Class14 | Class16 | Class28 | Class34 | Class35 | Class39 |
| True Labels |  | Class0           | 40     | 2      | 0       | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class0      |  | Class1           | 0      | 438    | 6       | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class1      |  | Class2           | 0      | 35     | 415     | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class2      |  | Class11          | 13     | 192    | 59      | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class11     |  | Class12          | 2      | 245    | 173     | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class12     |  | Class14          | 0      | 103    | 53      | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class14     |  | Class16          | 1      | 53     | 30      | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class16     |  | Class28          | 16     | 77     | 15      | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class28     |  | Class34          | 0      | 31     | 53      | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class34     |  | Class35          | 0      | 205    | 35      | 0       | 0       | 0       | 0       | 0       | 0       | 0       |
| Class35     |  | Class39          | 0      | 41     | 19      | 0       | 0       | 0       | 0       | 0       | 0       | 0       |

# Results & Conclusion

Given the results, it is clear that FGSM had the most impact (32.16% drop) on the baseline test accuracy, then PGD (22.03% drop) and then Gaussian Blur (9.69% drop).

Gaussian Blur:

- Deblurring and Contrast Enhancement did not significantly changed the test accuracy (fluctuates within ~1%).
- Canny Edge Concatenation did not significantly change the test accuracy (fluctuates within 3.22%)

FGSM:

- Denoising improved test accuracy by 10.97% & F1 score improved by around 0.1
- Contrast enhancement improved accuracy by 4.12% and F1 score improved by 0.12
- Canny Edge Concatenation decreased test accuracy by 22.17% & F1 score decreased by 0.16

PGD:

- Denoising improved test accuracy by 6.49% and F1 score improved by 0.05
- Contrast enhancement improved accuracy by 3.4% and F1 score improved by 0.07
- Canny Edge Concatenation improved accuracy by 5.37% and F1 score improved by 0.02
- Techniques such as color contrast, ensembling were not successful mitigation techniques

# Results & Conclusion

- Gaussian blur attacks are the most difficult to protect against at inference time
  - Black box attacks are more difficult to protect against
- Multiple techniques were able to improve accuracy on a PGD attack
  - Maybe it was a milder attack than FGSM (not entirely sure)
- The de-noising and contrast enhancement were effective against FGSM but canny edge concatenation was not
  - De-noising is directly combatting FGSM while canny edge is not
- Future works include:
  - Combining the successful mitigation techniques
  - Trying ensemble methods with features other than canny edge
  - Trying with a different dataset to ensure that results are consistent
  - Try mitigation techniques on various strengths of attacks
- It's difficult to improve performance on inference time alone, but not eager to rule it out

# Thank you!

Thank you to our Professors and TA mentor, Weiran, for all your help and guidance!

# Resources

- [1] Bassett, R. (2021). Color and Edge-Aware Adversarial Image Perturbations.
- [2] Hosseini, H., Xiao, B., & Poovendran, R. (2019). n-ML: Mitigating Adversarial Examples via Ensembles of Topologically Manipulated Classifiers.
- [3] Shafahi, A., Najibi, M., Ghiasi, A., Xu, Z., Dickerson, J., Studer, C., ... & Goldstein, T. (2019). Adversarial training for free!
- [4] Zamir S., Arora A., et al., (2021). CVPR. Multi-Stage Progressive Image Restoration.
- [5] Rath S. (2022), Traffic Sign Recognition using PyTorch and Deep Learning. Debugger Cafe.  
<https://debuggercafe.com/traffic-sign-recognition-using-pytorch-and-deep-learning/>
- [6] Color-Extraction-and-Image-Segmentation.(2020).  
<https://github.com/Ronny-22-Code/Color-Extraction-and-Image-Segmentation.git>

# Extra Slides

# Supporting materials (MobileNetV3)

```
baseline_original_images:Total number of test images: 3810  
baseline_original_images:Total correct predictions: 3783  
baseline_original_images:Accuracy: 99.291  
Average FPS: 7.297
```

```
baseline_gaussian_blur:Total number of test images: 3810  
baseline_gaussian_blur:Total correct predictions: 3414  
baseline_gaussian_blur:Accuracy: 89.606  
Average FPS: 8.765  
baseline_pgd_attack:Total number of test images: 3810  
baseline_pgd_attack:Total correct predictions: 2944  
baseline_pgd_attack:Accuracy: 77.270  
Average FPS: 8.290
```

```
baseline_fgsm_attack:Total number of test images: 3810  
baseline_fgsm_attack:Total correct predictions: 2558  
baseline_fgsm_attack:Accuracy: 67.139  
Average FPS: 8.629
```

```
de-gaussian_blur:Total number of test images: 3810  
de-gaussian_blur:Total correct predictions: 3343  
de-gaussian_blur:Accuracy: 87.743  
Average FPS: 7.327  
de-pgd_attack:Total number of test images: 3810  
de-pgd_attack:Total correct predictions: 3191  
de-pgd_attack:Accuracy: 83.753  
Average FPS: 7.455  
de-fgsm_attack:Total number of test images: 3810  
de-fgsm_attack:Total correct predictions: 2976  
de-fgsm_attack:Accuracy: 78.110  
Average FPS: 7.298
```

```
de-gaussian_blur_light:Total number of test images: 3810  
de-gaussian_blur_light:Total correct predictions: 3366  
de-gaussian_blur_light:Accuracy: 88.346  
Average FPS: 10.029
```

```
de-pgd_attack_light:Total number of test images: 3810  
de-pgd_attack_light:Total correct predictions: 3321  
de-pgd_attack_light:Accuracy: 87.165  
Average FPS: 9.774  
de-fgsm_attack_light:Total number of test images: 3810  
de-fgsm_attack_light:Total correct predictions: 3133  
de-fgsm_attack_light:Accuracy: 82.231  
Average FPS: 9.641
```

# Supporting materials (MobileNetV3)

When color averaging is applied

```
baseline_original_images_color_avg:Total number of test images: 3810
baseline_original_images_color_avg:Total correct predictions: 3653
baseline_original_images_color_avg:Accuracy: 95.879
Average FPS: 10.068
```

```
baseline_gaussian_blur_color_avg:Total number of test images: 3810
baseline_gaussian_blur_color_avg:Total correct predictions: 2777
baseline_gaussian_blur_color_avg:Accuracy: 72.887
Average FPS: 8.154
```

```
baseline_pgd_attack_color_avg:Total number of test images: 3810
baseline_pgd_attack_color_avg:Total correct predictions: 3004
baseline_pgd_attack_color_avg:Accuracy: 78.845
Average FPS: 7.634
```

```
baseline_fgsm_attack_color_avg:Total number of test images: 3810
baseline_fgsm_attack_color_avg:Total correct predictions: 2636
baseline_fgsm_attack_color_avg:Accuracy: 69.186
Average FPS: 8.139
```

```
de-gaussian_blur_color_avg:Total number of test images: 3810
de-gaussian_blur_color_avg:Total correct predictions: 3014
de-gaussian_blur_color_avg:Accuracy: 79.108
Average FPS: 7.603
```

```
de-pgd_attack_color_avg:Total number of test images: 3810
de-pgd_attack_color_avg:Total correct predictions: 2937
de-pgd_attack_color_avg:Accuracy: 77.087
Average FPS: 8.081
```

```
de-fgsm_attack_color_avg:Total number of test images: 3810
de-fgsm_attack_color_avg:Total correct predictions: 2696
de-fgsm_attack_color_avg:Accuracy: 70.761
Average FPS: 8.101
```

```
de-gaussian_blur_light_color_avg:Total number of test images: 3810
de-gaussian_blur_light_color_avg:Total correct predictions: 3061
de-gaussian_blur_light_color_avg:Accuracy: 80.341
Average FPS: 7.745
```

```
de-pgd_attack_light_color_avg:Total number of test images: 3810
de-pgd_attack_light_color_avg:Total correct predictions: 2910
de-pgd_attack_light_color_avg:Accuracy: 76.378
Average FPS: 9.964
```

```
de-fgsm_attack_light_color_avg:Total number of test images: 3810
de-fgsm_attack_light_color_avg:Total correct predictions: 2699
de-fgsm_attack_light_color_avg:Accuracy: 70.840
Average FPS: 10.125
```

# Supporting materials (MobileNetV3)

## F-1 scores

```
● (base) gman@Grants-MacBook-Air src % python f_1.py baseline_fgsm_attack  
[INFO]: Not loading pre-trained weights  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:20: DeprecationWarning: 'trained' is deprecated since 0.13 and may be removed in the future, please use 'weights'.  
warnings.warn(  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:22: DeprecationWarning: 'None' or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future, please use 'None'.  
warnings.warn('None', DeprecationWarning)  
[INFO]: Freezing hidden layers...  
F1 Score: 0.7540
```

```
(base) gman@Grants-Air src % python f_1.py baseline_pgd_attack_color_avg  
[INFO]: Not loading pre-trained weights  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:20: DeprecationWarning: 'trained' is deprecated since 0.13 and may be removed in the future, please use 'weights'.  
warnings.warn(  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:22: DeprecationWarning: 'None' or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future, please use 'None'.  
warnings.warn('None', DeprecationWarning)  
[INFO]: Freezing hidden layers...  
F1 Score: 0.8498
```

```
(base) gman@Grants-Air src % python f_1.py de-gaussian_blur_light  
[INFO]: Not loading pre-trained weights  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:20: DeprecationWarning: 'trained' is deprecated since 0.13 and may be removed in the future, please use 'weights'.  
warnings.warn(  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:22: DeprecationWarning: 'None' or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future, please use 'None'.  
warnings.warn('None', DeprecationWarning)  
[INFO]: Freezing hidden layers...  
F1 Score: 0.8836
```

```
● (base) gman@Grants-MacBook-Air src % python f_1.py baseline_fgsm_attack_color_avg  
[INFO]: Not loading pre-trained weights  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:20: DeprecationWarning: 'trained' is deprecated since 0.13 and may be removed in the future, please use 'weights'.  
warnings.warn(  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:22: DeprecationWarning: 'None' or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future, please use 'None'.  
warnings.warn('None', DeprecationWarning)  
[INFO]: Freezing hidden layers...  
F1 Score: 0.7677
```

```
● (base) gman@Grants-Air src % python f_1.py de-fgsm_attack  
[INFO]: Not loading pre-trained weights  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:20: DeprecationWarning: 'trained' is deprecated since 0.13 and may be removed in the future, please use 'weights'.  
warnings.warn(  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:22: DeprecationWarning: 'None' or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future, please use 'None'.  
warnings.warn('None', DeprecationWarning)  
[INFO]: Freezing hidden layers...  
F1 Score: 0.8566
```

```
● (base) gman@Grants-Air src % python f_1.py de-gaussian_blur_light_color_avg  
[INFO]: Not loading pre-trained weights  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:20: DeprecationWarning: 'trained' is deprecated since 0.13 and may be removed in the future, please use 'weights'.  
warnings.warn(  
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:22: DeprecationWarning: 'None' or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future, please use 'None'.  
warnings.warn('None', DeprecationWarning)  
[INFO]: Freezing hidden layers...  
F1 Score: 0.8412
```

# Supporting materials (MobileNetV3)

## F-1 scores

```
(base) gman@Grants-MacBook-Air src % python f_1.py baseline_gaussian_blur
[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated sinc
nd may be removed in the future. The current behavior is equivalent to passing 'weig
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.9196
```

```
F1 Score: 0.8566
(base) gman@Grants-Air src % python f_1.py de-fgsm_attack_color_avg
[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated sinc
nd may be removed in the future. The current behavior is equivalent to passing 'weig
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.7875
```

```
F1 Score: 0.8912
(base) gman@Grants-Air src % python f_1.py de-pgd_attack
[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated sinc
nd may be removed in the future. The current behavior is equivalent to passing 'weig
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.8953
```

```
F1 Score: 0.9350
(base) gman@Grants-MacBook-Air src % python f_1.py baseline_gaussian_blur_color_avg
[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated sinc
nd may be removed in the future. The current behavior is equivalent to passing 'weig
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.7766
```

```
F1 Score: 0.7073
(base) gman@Grants-Air src % python f_1.py de-fgsm_attack_light
[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated sinc
nd may be removed in the future. The current behavior is equivalent to passing 'weig
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.8737
```

```
F1 Score: 0.8355
(base) gman@Grants-Air src % python f_1.py de-pgd_attack_color_avg
[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated sinc
nd may be removed in the future. The current behavior is equivalent to passing 'weig
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.8384
```

# Supporting materials (MobileNetV3)

## F-1 scores

```
F1 Score: 0.7766
(base) gman@Grants-MacBook-Air src % python f_1.py baseline_original_images

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.9953
```

```
F1 Score: 0.9586
(base) gman@Grants-Air src % python f_1.py de-pgd_attack_light

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
[Open file in editor (cmd + click)]
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.9115
```

```
(base) gman@Grants-Air traffic_sign_color_smote % cd src
(base) gman@Grants-Air src % python f_1.py baseline_original_images_color_avg

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.9708
```

```
(base) gman@Grants-Air src % python f_1.py de-pgd_attack_light_color_avg

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.8178
```

```
F1 Score: 0.6757
(base) gman@Grants-Air src % python f_1.py de-fgsm_attack_light_color_avg

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.7719
```

```
(base) gman@Grants-MacBook-Air src % python f_1.py de-gaussian_blur

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.8998
```

```
(base) gman@Grants-Air src % python f_1.py de-pgd_attack_light_color_avg

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.8178
```

```
(base) gman@Grants-Air src % python f_1.py baseline_pgd_attack

[INFO]: Not loading pre-trained weights
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208:
ning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the
please use 'weights' instead.
    warnings.warn(
/Users/gman/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223:
ning: Arguments other than a weight enum or 'None' for 'weights' are deprecated s
nd may be removed in the future. The current behavior is equivalent to passing 'wei
'.
    warnings.warn(msg)
[INFO]: Freezing hidden layers...
F1 Score: 0.8431
```

# Solution Space (MobileNetV3)

Looking at only stop signs (113 images):

|  | Test Accuracy | F1 Score |
|--|---------------|----------|
| Baseline with FGSM Attack                    | 97.345%       | 0.986    |
| With Canny Edge and dimensionality reduction | 97.345%       | 0.986    |